LongMemEval Small development: native first outcomes

451 planned queries per method | preselected checkpoint | separate Reader / native Judge

A Missing-outcome bounds, all 451

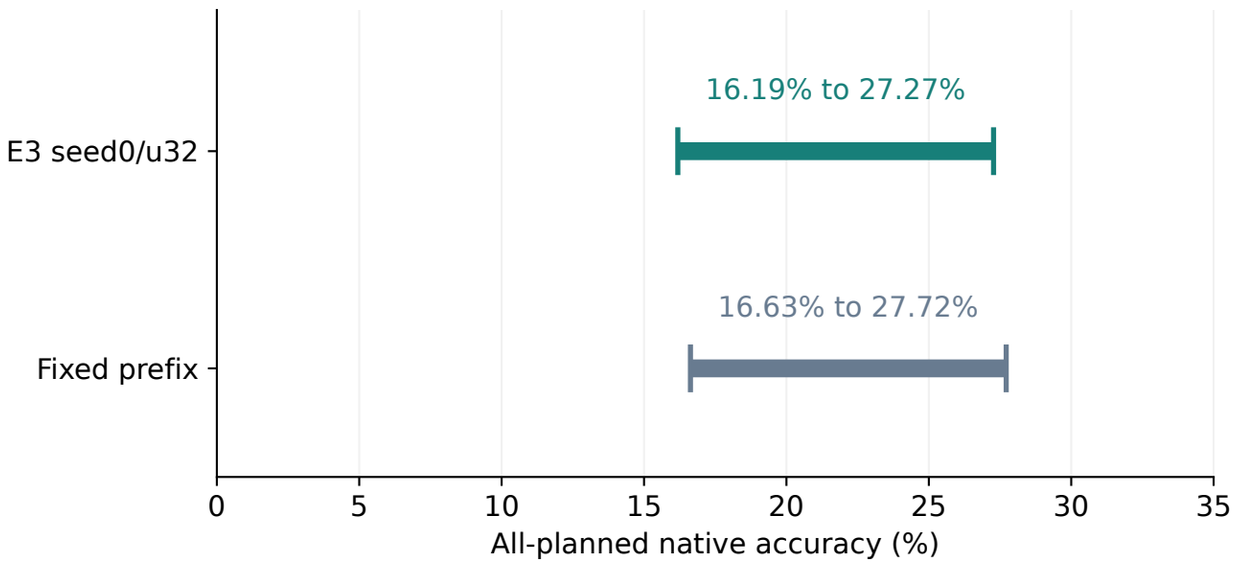

50 unscored per method; complete-cohort point metrics unavailable.
Intervals show missing-result bounds, not confidence intervals.

B Paired common-scored subset

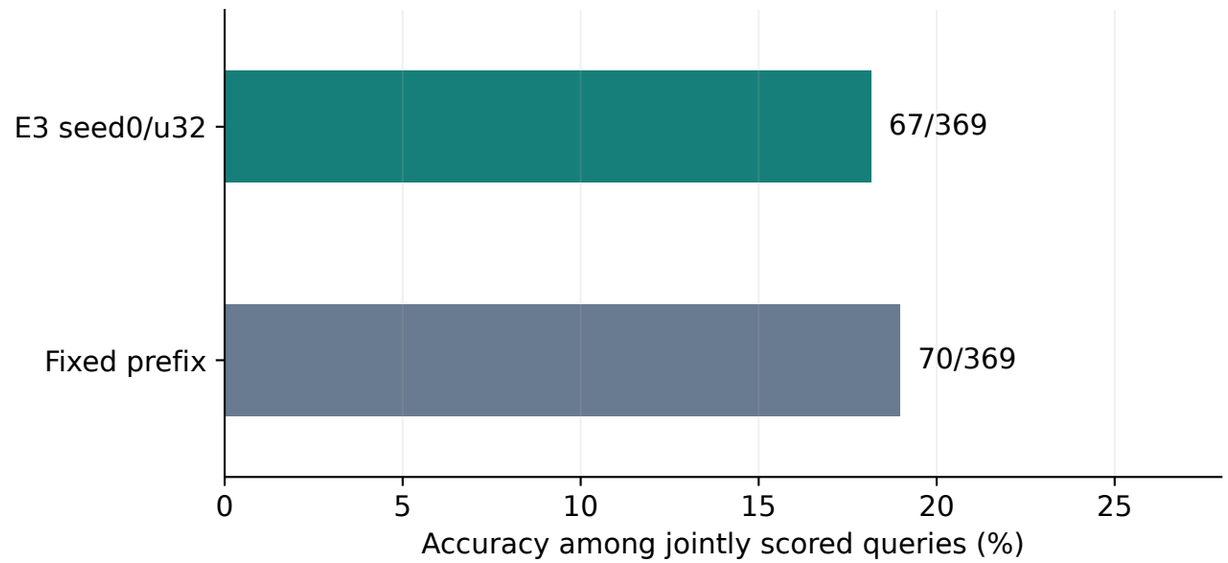

E3 gains 14 / losses 17 / ties 338; difference -0.81 pp.
Conditional on both completing; shared histories, no independent Source CI.

C READ selections and actual returns

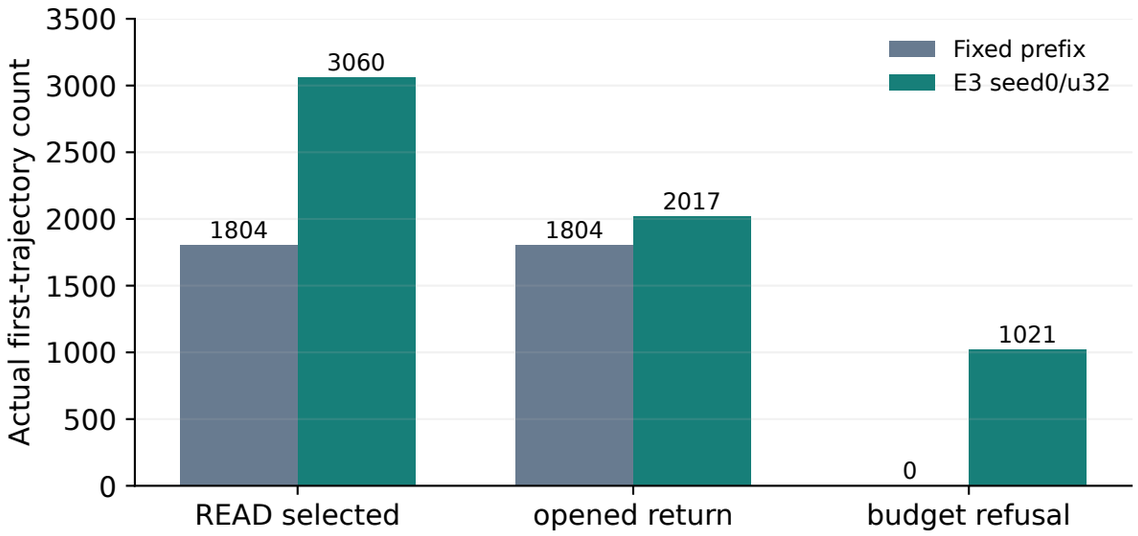

E3 also has 22 observed returns. Interrupted row9 ends in SEARCH.
449 complete E3 packets contain 1,885 original_source items.

D Prediction + scoring invocation costs

Measured quantity	Fixed prefix	E3 seed0/u32
Known HTTP tokens (M)	3.041	3.071
Whole client CPU (h)	10.590	8.986
Local NN calls	0	3,594
Local NN input tokens (M)	0.000	25.071
Model loads	0	51

CPU includes copy/capture/IO; nested CPU counted once.
Setup, audits, storage/archive, training and unmetered author work excluded.
GPU reservation intervals include waits; Host GPU busy time unmeasured.

Known Small development exposure; shared histories. First failures remain unscored. Source/input audit does not certify QA.
Primary physical step snapshots1-9 were never generated; raw evidence and metadata correction are disclosed in the report.